

Logistic Regression notes

Edward Hoang

December 23, 2021

Abstract

This is a Course note from week 2 of Neural Networks and Deep Learning Course

1 Neural Network notations

General comments:

- Superscript (i) will denote the i^{th} training example while superscript [l] will denote the l^{th} layer

Sizes:

- m: number of examples in the dataset.
- n_x : input size
- n_y : output size (or number of classes)
- $n_h^{[l]}$: number of hidden unit of the l^{th} layer

In a loop, it is possible to denote $n_x = n_h^{[0]}$ and $n_y = n_j^{[number\ of\ layers + 1]}$

- L: number of layers in the network.

Objects:

- $X \in R^{n_x \times m}$ is the input matrix
- $x^{(i)} \in R^{n_x}$ is the i^{th} example represented as a column vector
- $Y \in R^{n_y \times m}$ is the label matrix
- $y^i \in R^{n_y}$ is the output label for the i^{th} example
- $W^l \in R^{number\ of\ units\ in\ next\ layer \times numbers\ of\ units\ in\ the\ previous\ layer}$ is the weight matrix, superscript [l] indicates the layer
- $b^{[l]} \in R^{number\ of\ units\ in\ next\ layer}$ is the bias vector in the l^{th} layer
- $\hat{y} \in R^{n_y}$ is the predicted output vector. It can also be denoted $a^{[L]}$ where L is the number of layers in the network.

Common forward propagation equation examples

$a = g^{[l]}(W_x x^{(i)} + b_1) = g^{[l]}(z_1)$ where $g^{[l]}$ denotes the l^{th} layer activation function

$\hat{y}^{(i)} = softmax(W_h h + b_2)$

- General Activation Formula: $a_j^{[l]} = g^{[l]}(\sum_k w_{jk}^{[l-1]} a_k^{[l-1]} + b_k^{[l]}) = g^{[l]}(z_j^{[l]})$

- $J(x, W, b, y)$ or $J(\hat{y}, y)$ denote the cost function.

Examples of cost function:

- $J_{CE}(\hat{y}, y) = -\sum_{i=0}^m y(i) \log(\hat{y}^{(i)})$
- $J_1(\hat{y}, y) = \sum_{i=0}^m |y^{(i)} - \hat{y}^{(i)}|$

2 Logistic Regression notes

Personal notes:

- **Form:** $\sigma(z) = \frac{1}{1+e^{-z}}$ with $z = W^T x + b$
- **Dataset:** $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$
- **Want** $\hat{y}^{(i)} \approx y^{(i)}$
- **Need:**
 - + Lost function to evaluate how well the algorithm's doing for a single training example
 - + Form: $l(\hat{y}, y) = -(y \log \hat{y} + (1 - y) \log(1 - \hat{y}))$
 - + Cost function to evaluate how well the algorithm's doing for the entire training set
 - + Form: $Cost = \frac{1}{m} \sum_{i=1}^m l(\hat{y}, y)$
 - $or = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$
 - + Gradient Descent to take the model closer to the global minimum of the function
 - * **Learning rate:** how big a step we take on each iteration.
 - * **Form:**
 - Repeat {
 - $w := w - \alpha \frac{\partial J(w, b)}{\partial w}$
 - $b := b - \alpha \frac{\partial J(w, b)}{\partial b}$
 - }

Note: if J is a function of 2 or more variables, we use the symbol ∂ instead of "d" to denote the "partial derivative".